Adarsh Kumar Jha

ad4rush@gmail.com | +91-8287993780 | github.com/ad4rush | linkedin.com/in/ad4rush | ad4rush.github.io/PortFolio | Delhi, India

WORK EXPERIENCE

AI Engineer (Production Engineering) *Jan 2026 -- Present (Intern to Full-Time, Jun 2026)*
TBO.com (Travel Boutique Online)

- Architected near-real-time **ETL/ELT** lakehouse pipelines on S3 using **Apache Spark**, **AWS Glue**, **Kafka CDC (Debezium)**, and **Apache Hudi**; enabled sub-second incremental processing across **billions of records**.
- Streamlined **AWS Glue ETL pipeline reliability**, elevating job success rate from **94% to 99.8%** and cutting operating expenses by **$300/day** through resource tuning and schema caching.
- Restructured **distributed data processing** by routing **ECS** reads to same-AZ **Kafka brokers**, eliminating cross-AZ network latency and saving **$5,000+/month** in AWS transfer costs.
- Engineered a schema-aware **Text-to-SQL** ontology agent validating complex **SQL** on **Amazon Athena**, cutting **BI team workload by 35%**; exposed schema tools via **MCP** to Claude clients.
- Automated **LLM** classification of **650+ daily customer emails** and **100,000+ KAM notes** via **AWS Lambda**; applied prompt caching to slash inference spend by **70%+** and manual review by **57%**.
- Pioneered a **conversational AI** voice bot (**AWS Nova Sonic**, **RAG**, **tool calling**) for automated IVR triage, sustaining **p95 response latency under 600ms**.

ML Engineer Intern *May 2025 -- Oct 2025*
IIIT-Delhi -- Research to Commercialization (R2C) | *Advisor: Dr. Anuj Grover*

- Built a commercialization matchmaking platform with **FastAPI REST APIs** and **multi-agent AI** document parsing, advancing project toward **INR 40 Lakhs ($50K USD) NRDC funding**.
- Calibrated document intelligence models in **PyTorch** with **LoRA** and **PEFT**, reaching **97% extraction accuracy**, **75% parameter compression**, and **2x faster GPU inference throughput**.
- Engineered an asynchronous document queue via **Redis** and **Celery**, batch-parsing **500+ PDFs daily** with zero task failures and resilient automated retries across distributed worker pools.

Independent Research Project -- Ontology Path Extraction *Aug 2024 -- Dec 2024*
IIIT-Delhi x Oracle | Advisors: Prof. Raghava Mutharaju, Suman Roy (Oracle)

- Fine-tuned **R-BERT** in **PyTorch** with span-based **embeddings**, surpassing LLaMA-3 and GPT-3 baselines by **6 F1 points**; indexed **25,000+ biomedical entities** via FAISS, reducing false positives by **20%**.
- Benchmarked relation extraction against BioBERT, reaching **0.84 micro-F1** across 10 classes on NCBI datasets.

EDUCATION

Indraprastha Institute of Information Technology Delhi (IIIT-Delhi) *2022 -- 2026*
B.Tech, Computer Science and Artificial Intelligence *CGPA: 8.2 / 10*

PROJECTS

VICAS Lab Portal & Resume Knowledge Graph [Code] *Aug 2025 -- Apr 2026*
React, FastAPI, Docker, AWS EC2, Neo4j Knowledge Graphs, FAISS, GCN

- Deployed an AI talent marketplace on **React**, **FastAPI**, **Docker**, and **AWS EC2** across 6 microservices with OAuth2 and CI/CD; successfully onboarded **160+ registered students**.
- Formulated a **Gemini RAG** pipeline with **FAISS** semantic retrieval and **Neo4j knowledge graphs**; trained a **GCN** for skill link prediction with cross-encoder re-ranking scoring **91.4% precision**.

ML Serving Diagnostics & Observability Platform [Code] *Jan 2025 -- Apr 2025*
Python, FastAPI, Docker, Kubernetes, TensorFlow, PyTorch, Google Cloud Storage

- Constructed an automated diagnostics engine for **Kubernetes** ML workloads, isolating root causes of OOM faults and latency spikes; slashed mean time to resolution by **40%** across 5 production services.
- Implemented a **FastAPI** telemetry dashboard streaming node health to **Google Cloud Storage**, cataloging proactive remediation runbooks that prevented service downtime.

Customer Churn Prediction -- Ensemble ML Benchmarking [Code] *Aug 2024 -- Dec 2024*
Python, Scikit-Learn, XGBoost, LightGBM, CatBoost, Optuna

- Benchmarked **XGBoost, LightGBM, and CatBoost** on 50,000+ customer profiles, optimizing hyperparameters with Optuna to achieve **89.4% ROC-AUC** (a **12% gain** over baseline trees).
- Evaluated precision-recall curves to minimize false negatives, saving simulated revenue loss by **18%** across test cohorts.

TECHNICAL SKILLS

Languages: Python, SQL (Complex Queries, Optimization, Window Functions), Java, C, C++, JavaScript
Data Systems: Apache Spark, Apache Kafka, Debezium (CDC), AWS Glue, Apache Hudi, Parquet, Amazon Athena, ETL/ELT, Data Lakehouse, DynamoDB, Neo4j, Distributed Stream Processing
AI / Machine Learning: PyTorch, LLMs, RAG, AI Agents, Tool Calling, MCP, Text-to-SQL, FAISS Vector Search, Embeddings, GCN, LoRA, PEFT, XGBoost, LightGBM, CatBoost
Cloud & DevOps: AWS (Athena, Glue, S3, Lambda, ECS, DynamoDB, CloudWatch, EC2), GCP (BigQuery), FastAPI, Docker, Kubernetes, Linux, Git, CI/CD

ACHIEVEMENTS & LEADERSHIP

Ranked AIR 252 in GATE 2026 (Computer Science & Information Technology) *2026*
JEE Mains 2022: All India Rank **4778 (99.5 percentile)** among 872,970 candidates *2022*
Codeforces (Specialist - Max Rating 1450+), LeetCode 300+ problems solved, GFG *2023 -- 2026*
Operations Head, Odyssey IIIT-D -- Led 30-member logistics council for cultural festival hosting 30,000+ attendees *2025*